

AI ASSISTED CODING

ASSIGNMENT-17.3

Name: MD AMAN

Roll No: 2303A510C0

Batch:06

Task 1 Employee Dataset Cleaning

Task: Use AI assistance to clean an employee information dataset.

Prompt:

T ****Prompt:**** Clean an employee dataset using Python (Pandas & Scikit-learn): handle missing values (salary → median, designation → mode, experience → median), remove duplicate employee_id, convert text columns to lowercase, and encode designation (label encoding) and employment_type (one-hot encoding). Generate clean, complete code.

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

# Create a sample DataFrame for demonstration purposes
data = {
    'employee_id': [101, 102, 103, 104, 105, 101, 106, 107, 108, 109],
    'name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Alice', 'Frank', 'Grace', 'Heidi', 'Ivan'],
    'department': ['HR', 'IT', 'Finance', 'IT', 'Marketing', 'HR', 'IT', 'Finance', 'HR', 'Marketing'],
    'designation': ['Manager', 'Developer', 'Analyst', 'Developer', 'Marketing Specialist', 'Manager', 'Developer', np.nan, 'HR Assist', 'Developer'],
    'salary': [70000, 85000, 60000, np.nan, 75000, 70000, 90000, 62000, np.nan, 68000],
    'experience': [5, 8, 3, 7, np.nan, 5, 9, 4, 2, 6],
    'employment_type': ['Full-Time', 'Contract', 'Full-Time', 'Part-Time', 'Full-Time', 'Full-Time', 'Contract', 'Full-Time', 'Part-Time', 'Full-Time']
}
df = pd.DataFrame(data)

print("Original DataFrame:")
display(df.head(10))
print("\nOriginal DataFrame Info:")
df.info()
```

... Original DataFrame:

	employee_id	name	department	designation	salary	experience	employment_type
0	101	Alice	HR	Manager	70000.0	5.0	Full-Time
1	102	Bob	IT	Developer	85000.0	8.0	Contract
2	103	Charlie	Finance	Analyst	60000.0	3.0	Full-Time
3	104	David	IT	Developer	NaN	7.0	Part-Time
4	105	Eve	Marketing	Marketing Specialist	75000.0	NaN	Full-Time
5	101	Alice	HR	Manager	70000.0	5.0	Full-Time
6	106	Frank	IT	Developer	90000.0	9.0	Contract
7	107	Grace	Finance	NaN	62000.0	4.0	Full-Time
8	108	Heidi	HR	HR Assistant	NaN	2.0	Part-Time

Code:

```

import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

# 2. Handle Missing Values
df['salary'].fillna(df['salary'].median(), inplace=True)
df['designation'].fillna(df['designation'].mode()[0], inplace=True)
df['experience'].fillna(df['experience'].median(), inplace=True)
print("\nDataFrame Info after handling missing values:")
df.info()

# 3. Remove Duplicate Employee IDs
df.drop_duplicates(subset='employee_id', inplace=True)
print("\nDataFrame Info after removing duplicate 'employee_id':")
df.info()

# 4. Convert Text Columns to Lowercase
text_columns = ['name', 'department', 'designation', 'employment_type']
for col in text_columns:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = df[col].str.lower()

# 5. Encode Categorical Features
# Label Encoding for 'designation'
le = LabelEncoder()
df['designation_encoded'] = le.fit_transform(df['designation'])

# One-Hot Encoding for 'employment_type'
df['employment_type'].fillna('unknown', inplace=True)
ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
employment_type_encoded = ohe.fit_transform(df[['employment_type']])
employment_type_df = pd.DataFrame(employment_type_encoded, columns=ohe.get_feature_names_out(['employment_type']), index=df.index)
df = pd.concat([df, employment_type_df], axis=1)

# Drop original categorical columns
df.drop(columns=['designation', 'employment_type'], inplace=True)

```

Output:

```

df['employment_type'].fillna('unknown', inplace=True) # or choose another strategy

```

	employee_id	name	department	salary	experience	designation_encoded	employment_type_contract	employment_type_full-time
0	101	alice	hr	70000.0	5.0	3	0.0	1.0
1	102	bob	it	85000.0	8.0	1	1.0	0.0
2	103	charlie	finance	60000.0	3.0	0	0.0	1.0
3	104	david	it	70000.0	7.0	1	0.0	0.0
4	105	eve	marketing	75000.0	5.0	4	0.0	1.0
6	106	frank	it	90000.0	9.0	1	1.0	0.0
7	107	grace	finance	62000.0	4.0	1	0.0	1.0
8	108	heidi	hr	70000.0	2.0	2	0.0	0.0
9	109	ivan	marketing	68000.0	6.0	0	0.0	0.0

```

Cleaned DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
Index: 9 entries, 0 to 9
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   employee_id                           9 non-null     int64
1   name                                  9 non-null     object
2   department                            9 non-null     object
3   salary                                9 non-null     float64
4   experience                            9 non-null     float64
5   designation_encoded                    9 non-null     int64
6   employment_type_contract               9 non-null     float64
7   employment_type_full-time              9 non-null     float64
8   employment_type_part-time              9 non-null     float64
9   employment_type_unknown                9 non-null     float64
dtypes: float64(6), int64(2), object(2)
memory usage: 792.0+ bytes

```

Task 2 – Student Performance Data Preprocessing Task:

Preprocess student performance data using AI assistance.

Instructions:

- Convert exam_date column into datetime format
- Extract new features (month, semester, day)
- Normalize total_marks
- Detect and handle outliers in marks using IQR or Z-score

Prompt:

Preprocess a student performance dataset using Python (Pandas & Scikit-learn): convert exam_date to datetime, extract month, semester, and day features, normalize total_marks, and detect & handle outliers using IQR or Z-score. Generate clean and complete code with the transformed dataset as output.

T Preprocess a student performance dataset using Python (Pandas & Scikit-learn): convert exam_date to datetime, extract month, semester, and day features, normalize total_marks, and detect & handle outliers using IQR or Z-score. Generate clean and complete code with the transformed dataset as output.

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
# Create a sample DataFrame for demonstration
data = {
    'student_id': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15],
    'exam_date': [
        '2023-01-15', '2023-03-20', '2023-05-10', '2023-09-01', '2023-11-25',
        '2023-01-22', '2023-03-11', '2023-05-19', '2023-09-15', '2023-11-05',
        '2023-02-01', '2023-04-01', '2023-06-01', '2023-10-01', '2023-12-01'
    ],
    'course': [
        'Math', 'Science', 'Math', 'History', 'Science',
        'Math', 'Science', 'Math', 'History', 'Science',
        'Math', 'Science', 'Math', 'History', 'Science'
    ],
    'attendance_rate': [
        0.95, 0.88, 0.75, 0.92, 0.80, 0.99, 0.85, 0.70, 0.90, 0.82,
        0.93, 0.87, 0.78, 0.91, 0.81
    ],
    'total_marks': [
        85, 72, 60, 91, 78, 98, 65, 55, 88, 70,
        80, 75, 68, 95, 77
    ]
}
df = pd.DataFrame(data)

print("Original DataFrame:")
display(df.head())
print("\nOriginal DataFrame Info:")
df.info()
```

```
*** Original DataFrame:
  student_id  exam_date  course  attendance_rate  total_marks
0          1  2023-01-15   Math              0.95           85
1          2  2023-03-20  Science              0.88           72
2          3  2023-05-10   Math              0.75           60
3          4  2023-09-01  History              0.92           91
4          5  2023-11-25  Science              0.80           78

Original DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15 entries, 0 to 14
Data columns (total 5 columns):
#   Column              Non-Null Count  Dtype
---
```

Code:

```

import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Assuming 'df' DataFrame is already loaded from the previous cell.

# 1. Convert exam_date to Datetime and Extract Features
df['exam_date'] = pd.to_datetime(df['exam_date'])

df['exam_month'] = df['exam_date'].dt.month
df['exam_day_of_week'] = df['exam_date'].dt.dayofweek # Monday=0, Sunday=6

def get_semester(month):
    if 1 <= month <= 5:
        return 1 # Semester 1: Jan - May
    elif 6 <= month <= 8:
        return 2 # Semester 2: Jun - Aug
    else:
        return 3 # Semester 3: Sep - Dec

df['semester'] = df['exam_month'].apply(get_semester)

# 2. Normalize total_marks
scaler = MinMaxScaler()
df['total_marks_normalized'] = scaler.fit_transform(df[['total_marks']])

# 3. Detect and Handle Outliers in total_marks using IQR
Q1 = df['total_marks'].quantile(0.25)
Q3 = df['total_marks'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Handle outliers by capping (winsorization)
df['total_marks_handled_outliers'] = df['total_marks'].clip(lower=lower_bound, upper=upper_bound)

print("Final Transformed DataFrame:")
display(df.head(10))
print("\nFinal Transformed DataFrame Info:")
df.info()

```

```

# Label Encoding for 'designation'
le = LabelEncoder()
df['designation_encoded'] = le.fit_transform(df['designation'])
print(f"Designation mappings: {list(le.classes_)}")

# One-Hot Encoding for 'employment_type'
# First, handle potential missing values in 'employment_type' for OneHotEncoder compatibility
df['employment_type'].fillna('unknown', inplace=True) # or choose another strategy
ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
employment_type_encoded = ohe.fit_transform(df[['employment_type']])

# Create a DataFrame from the one-hot encoded array with appropriate column names
employment_type_df = pd.DataFrame(employment_type_encoded, columns=ohe.get_feature_names_out(['employment_type']), index=df.index)

# Concatenate the new one-hot encoded columns with the original DataFrame
df = pd.concat([df, employment_type_df], axis=1)

# Optionally, drop the original 'designation' and 'employment_type' columns if you only want the encoded ones
df.drop(columns=['designation', 'employment_type'], inplace=True)

print("\nCleaned DataFrame after encoding categorical features:")
display(df.head(10))
print("\nCleaned DataFrame Info:")
df.info()

```

Output:

```
''' Designation mappings: ['analyst', 'developer', 'hr assistant', 'manager', 'marketing specialist']

Cleaned DataFrame after encoding categorical features:
/tmp/ipykernel_3611/429617752.py:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the
original object.

df['employment_type'].fillna('unknown', inplace=True) # or choose another strategy
employee_id  name  department  salary  experience  designation_encoded  employment_type_contract  employment_type_full-time  employment_type_part-time  employment_type_unknown
0          101    alice        hr  70000.0        5.0            3              0.0              1.0              0.0              0.0
1          102     bob        it  85000.0        8.0            1              1.0              0.0              0.0              0.0
2          103  charlie    finance  60000.0        3.0            0              0.0              1.0              0.0              0.0
3          104   david        it  70000.0        7.0            1              0.0              0.0              1.0              0.0
4          105    eve    marketing  75000.0        5.0            4              0.0              1.0              0.0              0.0
6          106   frank        it  90000.0        9.0            1              1.0              0.0              0.0              0.0
7          107   grace    finance  62000.0        4.0            1              0.0              1.0              0.0              0.0
8          108    heidi        hr  70000.0        2.0            2              0.0              0.0              1.0              0.0
9          109    ivan    marketing  68000.0        6.0            0              0.0              0.0              0.0              1.0

Cleaned DataFrame Info:
(class 'pandas.core.frame.DataFrame')
Index: 9 entries, 0 to 9
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype
---  --
0   employee_id         9 non-null     int64
1   name                9 non-null     object
2   department          9 non-null     object
3   salary              9 non-null     float64
4   experience           9 non-null     float64
5   designation_encoded  9 non-null     int64
6   employment_type_contract  9 non-null     float64
7   employment_type_full-time  9 non-null     float64
8   employment_type_part-time  9 non-null     float64
9   employment_type_unknown  9 non-null     float64
dtypes: float64(8), int64(2), object(2)
memory usage: 792.0+ bytes
```

Task 3 – Student Health Data Preparation

Task: Clean and preprocess student health records using AI scripts.

Instructions:


- Handle missing values in height, weight, BMI
- Encode categorical features (medical_condition, fitness_status)
- Scale numerical features using standard scaling
- Split dataset into training and testing sets

Expected Output:

A preprocessed dataset ready for predictive health analysis.

Prompt:

Preprocess a student health dataset using Python (Pandas & Scikit-learn): handle missing values (height, weight, BMI), encode categorical features (medical_condition, fitness_status), apply standard scaling to numerical features, and split the dataset into training and testing sets. Generate clean and complete code with the final processed output. create a separate code for creating the dataset and give me all required codes in only one chat and finally output



Dataset:

```
import pandas as pd
import numpy as np

# Create a sample DataFrame for demonstration of student health data
data = {
    'student_id': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'height_cm': [165, 172, 158, 180, 170, 163, np.nan, 175, 168, 171],
    'weight_kg': [60, 75, 55, np.nan, 68, 62, 70, 78, 65, np.nan],
    'BMI': [22.0, 25.5, 21.9, 24.7, np.nan, 23.3, 26.5, 25.1, 22.9, 23.9],
    'medical_condition': ['None', 'Asthma', 'Allergies', 'None', 'None', 'Asthma', 'None', 'Diabetes', 'Allergies', 'None'],
    'fitness_status': ['Good', 'Average', 'Poor', 'Good', 'Average', 'Good', 'Poor', 'Average', 'Good', 'Average']
}

df_health = pd.DataFrame(data)
print("Original Student Health DataFrame:")
display(df_health.head(10))
print("\nOriginal Student Health DataFrame Info:")
df_health.info()
```

Original Student Health DataFrame:

	student_id	height_cm	weight_kg	BMI	medical_condition	fitness_status
0	1	165.0	60.0	22.0	None	Good
1	2	172.0	75.0	25.5	Asthma	Average
2	3	158.0	55.0	21.9	Allergies	Poor
3	4	180.0	NaN	24.7	None	Good
4	5	170.0	68.0	NaN	None	Average
5	6	163.0	62.0	23.3	Asthma	Good
6	7	NaN	70.0	26.5	None	Poor
7	8	175.0	78.0	25.1	Diabetes	Average
8	9	168.0	65.0	22.9	Allergies	Good
9	10	171.0	NaN	23.9	None	Average

Original Student Health DataFrame Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   student_id            10 non-null    int64
1   height_cm             9 non-null     float64
2   weight_kg             8 non-null     float64
3   BMI                   9 non-null     float64
4   medical_condition     10 non-null    object
5   fitness_status        10 non-null    object
dtypes: float64(3), int64(1), object(2)
```

Code:

```
import numpy as np
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer

# Assuming df_health is already loaded from the previous cell
# Make a copy to avoid modifying the original DataFrame if it's reused later
df_processed = df_health.copy()

# 1. Handle Missing Values
# Using median for numerical columns (height, weight, BMI)
for col in ['height_cm', 'weight_kg', 'BMI']:
    if df_processed[col].isnull().any():
        median_val = df_processed[col].median()
        df_processed[col].fillna(median_val, inplace=True)

print("DataFrame after handling missing values:")
display(df_processed.head())

# 2. Encode Categorical Features
# One-Hot Encoding for 'medical_condition' and 'fitness_status'
# For medical_condition, one-hot is appropriate as there's no inherent order.
# For fitness_status, one-hot is also used for simplicity, assuming no strict ordinality desired.
ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
# Encoding 'medical_condition'
medical_encoded = ohe.fit_transform(df_processed[['medical_condition']])
medical_df = pd.DataFrame(medical_encoded, columns=ohe.get_feature_names_out(['medical_condition']), index=df_processed.index)
df_processed = pd.concat([df_processed, medical_df], axis=1)
# Encoding 'fitness_status'
fitness_encoded = ohe.fit_transform(df_processed[['fitness_status']])
fitness_df = pd.DataFrame(fitness_encoded, columns=ohe.get_feature_names_out(['fitness_status']), index=df_processed.index)
df_processed = pd.concat([df_processed, fitness_df], axis=1)

# Drop original categorical columns
df_processed.drop(columns=['medical_condition', 'fitness_status'], inplace=True)
print("\nDataFrame after encoding categorical features:")
display(df_processed.head())

# 3. Apply Standard Scaling to Numerical Features
numerical_cols = ['height_cm', 'weight_kg', 'BMI']
scaler = StandardScaler()
df_processed[numerical_cols] = scaler.fit_transform(df_processed[numerical_cols])
print("\nDataFrame after applying Standard Scaling:")
display(df_processed.head())

# 4. Split the dataset into training and testing sets
# Assuming 'student_id' is just an identifier and not a feature, and no specific target column is given.
# For demonstration, let's treat all processed columns except 'student_id' as features (X).
# If there was a target variable (e.g., 'performance'), it would be separated.
X = df_processed.drop(columns=['student_id'])
# Since no explicit target variable (y) is specified, we'll just split X.
# In a real scenario, you would have X and y.
X_train, X_test = train_test_split(X, test_size=0.2, random_state=42)
print("\nShape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("\nFinal Processed Training Data (first 5 rows):")
```

Output:

```
Final Transformed DataFrame:
```

	student_id	exam_date	course	attendance_rate	total_marks	exam_month	exam_day_of_week	semester	total_marks_normalized	total_marks_handled_outliers
0	1	2023-01-15	Math	0.95	85	1	6	1	0.697674	85
1	2	2023-03-20	Science	0.88	72	3	0	1	0.395349	72
2	3	2023-05-10	Math	0.75	60	5	2	1	0.116279	60
3	4	2023-09-01	History	0.92	91	9	4	3	0.837209	91
4	5	2023-11-25	Science	0.80	78	11	5	3	0.534884	78
5	6	2023-01-22	Math	0.99	98	1	6	1	1.000000	98
6	7	2023-03-11	Science	0.85	65	3	5	1	0.232558	65
7	8	2023-05-19	Math	0.70	55	5	4	1	0.000000	55
8	9	2023-09-15	History	0.90	88	9	4	3	0.767442	88
9	10	2023-11-05	Science	0.82	70	11	6	3	0.348837	70

```
Final Transformed DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15 entries, 0 to 14
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   student_id            15 non-null    int64
1   exam_date             15 non-null    datetime64[ns]
2   course                15 non-null    object
3   attendance_rate        15 non-null    float64
4   total_marks           15 non-null    int64
5   exam_month            15 non-null    int32
6   exam_day_of_week      15 non-null    int32
7   semester              15 non-null    int64
8   total_marks_normalized 15 non-null    float64
9   total_marks_handled_outliers 15 non-null    int64
dtypes: datetime64[ns](1), float64(2), int32(2), int64(4), object(1)
memory usage: 1.2+ KB
```

Task 4: Real-Time Application: Ride-Sharing Data Cleaning

Scenario: A ride-sharing company has trip data with errors.

Requirements:

- Standardize pickup and drop locations
 - Fill missing fare amounts with median values
 - Remove duplicate ride entries
 - Normalize driver ratings
 - Generate a summary of data improvements
- Expected Output:

A clean dataset ready for operational analysis.

Prompt:

Clean a ride-sharing dataset using Python (Pandas & Scikit-learn): standardize pickup and drop locations, fill missing fare values with median, remove duplicate records, normalize driver ratings, and generate a summary of data cleaning steps. Produce clean, complete code with the final processed dataset. create a separately code for creating the adatset and give me all required codes in onluy one chat and finally output

T Clean a ride-sharing dataset using Python (Pandas & Scikit-learn): standardize pickup and drop locations, fill missing fare values with median, remove duplicate records, normalize driver ratings, and generate a summary of data cleaning steps. Produce clean, complete code with the final processed dataset. create a separately code for creating the adatset and give me all required codes in onluy one chat and finally output

DataSet:

```

import pandas as pd
import numpy as np
# Create a sample DataFrame for demonstration of ride-sharing data
data = {
    'ride_id': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12],
    'pickup_location': [
        'Downtown', 'Airport', 'Suburb', 'Downtown', 'Airport',
        'Suburb', 'Downtown', 'Airport', 'Downtown', 'Suburb',
        'Downtown', 'Airport', 'Downtown'
    ],
    'dropoff_location': [
        'Suburb', 'Downtown', 'Airport', 'Suburb', 'Downtown',
        'Airport', 'Suburb', 'Downtown', 'Airport', 'Downtown',
        'Suburb', 'Downtown', 'Suburb'
    ],
    'fare': [15.50, 25.00, 12.75, 16.00, np.nan, 13.00, 15.50, 26.00, 14.50, np.nan, 15.50, 24.00, 16.00],
    'distance_miles': [5.2, 12.1, 4.5, 6.0, 10.5, 4.8, 5.2, 12.5, 5.8, 4.0, 5.2, 11.9, 6.1],
    'driver_rating': [4.8, 3.5, 4.9, 4.2, 3.9, 5.0, 4.8, 4.1, 4.7, 3.8, 4.8, 4.0, 4.3],
    'ride_date': pd.to_datetime([
        '2023-01-01', '2023-01-01', '2023-01-02', '2023-01-02', '2023-01-03',
        '2023-01-03', '2023-01-04', '2023-01-04', '2023-01-05', '2023-01-05',
        '2023-01-01', '2023-01-06', '2023-01-06'
    ])
}
df_rides = pd.DataFrame(data)

print("Original Ride-Sharing DataFrame:")
display(df_rides.head(10))
print("\nOriginal Ride-Sharing DataFrame Info:")
df_rides.info()

```

```

Original Ride-Sharing DataFrame:

```

	ride_id	pickup_location	dropoff_location	fare	distance_miles	driver_rating	ride_date
0	1	Downtown	Suburb	15.50	5.2	4.8	2023-01-01
1	2	Airport	Downtown	25.00	12.1	3.5	2023-01-01
2	3	Suburb	Airport	12.75	4.5	4.9	2023-01-02
3	4	Downtown	Suburb	16.00	6.0	4.2	2023-01-02
4	5	Airport	Downtown	NaN	10.5	3.9	2023-01-03
5	6	suburb	Airport	13.00	4.8	5.0	2023-01-03
6	7	Downtown	suburb	15.50	5.2	4.8	2023-01-04
7	8	Airport	Downtown	26.00	12.5	4.1	2023-01-04
8	9	Downtown	Airport	14.50	5.8	4.7	2023-01-05

Code:

```

import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
# Assuming df_rides is already loaded from the previous cell
# Make a copy to avoid modifying the original DataFrame if it's reused later
df_processed_rides = df_rides.copy()
# Store initial and current state for summary
initial_rows = len(df_processed_rides)
initial_missing_fares = df_processed_rides['fare'].isnull().sum()
# 1. Standardize pickup and drop locations
# Convert to lowercase and strip whitespace
df_processed_rides['pickup_location'] = df_processed_rides['pickup_location'].str.lower().str.strip()
df_processed_rides['dropoff_location'] = df_processed_rides['dropoff_location'].str.lower().str.strip()
print("Step 1: Standardized 'pickup_location' and 'dropoff_location' to lowercase and stripped whitespace.")
# 2. Fill missing fare values with median
median_fare = df_processed_rides['fare'].median()
df_processed_rides['fare'].fillna(median_fare, inplace=True)
print(f"Step 2: Filled {initial_missing_fares} missing 'fare' values with the median ({median_fare}).")
# 3. Remove duplicate records based on all columns
df_processed_rides.drop_duplicates(inplace=True)
duplicated_rows_removed = initial_rows - len(df_processed_rides)
print(f"Step 3: Removed {duplicated_rows_removed} duplicate records.")
# 4. Normalize driver ratings using Min-Max Scaling (scale to 0-1 range)
scaler = MinMaxScaler()
df_processed_rides['driver_rating_normalized'] = scaler.fit_transform(df_processed_rides[['driver_rating']])
print("Step 4: Normalized 'driver_rating' using Min-Max Scaling, creating 'driver_rating_normalized' column.")
print("\n--- Data Cleaning Summary ---")
print(f"Initial number of rows: {initial_rows}")
print(f"Rows after removing duplicates: {len(df_processed_rides)}")
print(f"Missing 'fare' values initially: {initial_missing_fares}")
print(f"Missing 'fare' values after treatment: {df_processed_rides['fare'].isnull().sum()}")
print(f"Locations standardized (lowercase, strip whitespace): 'pickup_location', 'dropoff_location'")
print(f"Driver ratings normalized (Min-Max Scaling): 'driver_rating_normalized'")
print("\nFinal Processed Ride-Sharing DataFrame (first 10 rows):")
display(df_processed_rides.head(10))
print("\nFinal Processed Ride-Sharing DataFrame Info:")
df_processed_rides.info()

```

Output:

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].meth

df_processed_rides['fare'].fillna(median_fare, inplace=True)

	ride_id	pickup_location	dropoff_location	fare	distance_miles	driver_rating	ride_date	driver_rating_normalized
0	1	downtown	suburb	15.50	5.2	4.8	2023-01-01	0.866667
1	2	airport	downtown	25.00	12.1	3.5	2023-01-01	0.000000
2	3	suburb	airport	12.75	4.5	4.9	2023-01-02	0.933333
3	4	downtown	suburb	16.00	6.0	4.2	2023-01-02	0.466667
4	5	airport	downtown	15.50	10.5	3.9	2023-01-03	0.266667
5	6	suburb	airport	13.00	4.8	5.0	2023-01-03	1.000000
6	7	downtown	suburb	15.50	5.2	4.8	2023-01-04	0.866667
7	8	airport	downtown	26.00	12.5	4.1	2023-01-04	0.400000
8	9	downtown	airport	14.50	5.8	4.7	2023-01-05	0.800000
9	10	suburb	downtown	15.50	4.0	3.8	2023-01-05	0.200000

Final Processed Ride-Sharing DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
Index: 12 entries, 0 to 12
Data columns (total 8 columns):
Column Non-Null Count Dtype

0 ride_id 12 non-null int64
1 pickup_location 12 non-null object
2 dropoff_location 12 non-null object
3 fare 12 non-null float64
4 distance_miles 12 non-null float64
5 driver_rating 12 non-null float64
6 ride_date 12 non-null datetime64[ns]
7 driver_rating_normalized 12 non-null float64
dtypes: datetime64[ns](1), float64(4), int64(1), object(2)
memory usage: 864.0+ bytes